

Deep Q-Networks on Atari

Reinforcement Learning Homework — BIU 2026

Aviv Weidenfeld Shani Or Levy

May 2026

Contents

1	Game Choice	2
2	Preprocessing	2
3	Network Architecture	2
4	Final Hyperparameters	3
5	Training Protocol	3
6	Training Curves	4
7	Final Results	5
8	What Worked	5
8.1	Life-loss terminal signal	5
8.2	Fast epsilon annealing (250k steps)	5
8.3	Lower discount factor ($\gamma = 0.97$)	5
8.4	LR warmdown (<code>lr_decay_start_step = 3M</code>)	5
8.5	Best-checkpoint evaluation (<code>use_best_checkpoint_eval</code>)	6
8.6	Stable learning-rate regime	6
9	What Did Not Work	6
9.1	The post-peak collapse problem	6
10	Implementation Notes	7
11	Conclusion	7

1 Game Choice

We trained the DQN agent on **Alien** (ALE/Alien-v5).

Why Alien?

Alien is a moderately complex Atari game with a well-defined reward signal (shooting aliens, collecting power pellets, progressing through mazes). It has a meaningful gap between random and trained agent performance, making it an informative benchmark.

Threshold	Score	Status
Random agent (minimum must)	227.8	✓ Passed
Good grade (best linear agent)	950.0	✓ Passed — avg 3170.6
Bonus grade (DQN paper score)	3000.0	✓ Passed — avg 3170.6

2 Preprocessing

We follow the preprocessing from Mnih et al. (2013) closely:

1. **Grayscale conversion:** Each RGB frame ($210 \times 160 \times 3$) is converted to grayscale using `cv2.COLOR_RGB2GRAY`.
2. **Resize and crop:** The grayscale frame is resized to 84×110 using `cv2.INTER_AREA` interpolation, then cropped to 84×84 by removing 18 pixels from the top and 8 from the bottom (rows 18–102).
3. **Frame skipping:** The agent repeats each selected action for 4 raw environment steps (`frame_skip=4`), accumulating reward across them. This matches the paper’s action repetition.
4. **Max-pooling over last 2 frames:** To eliminate sprite flickering (a known ALE artefact), the observation fed to the network is the pixel-wise maximum of the last 2 frames within each skip block.
5. **Frame stacking:** The last 4 preprocessed frames are stacked into a $(4 \times 84 \times 84)$ tensor, giving the network temporal context. Frames are stored as `uint8` in the replay buffer and cast to `float32` in $[0, 1]$ at training time.
6. **Reward clipping:** Rewards are clipped to $\{-1, 0, +1\}$ via `np.sign(reward)`, as in the paper.
7. **Life-loss terminal:** Each life loss during a game is treated as a terminal state (the episode reset is not triggered, only the Q-value bootstrap chain is cut). This teaches the agent to avoid dying.

3 Network Architecture

We use the convolutional architecture from the 2013 DQN paper:

Layer	Type	Filters/Units	Kernel/Stride	Activation
Conv1	Conv2D	16	8×8 , stride 4	ReLU
Conv2	Conv2D	32	4×4 , stride 2	ReLU
FC1	Linear	256	—	ReLU
Output	Linear	$ \mathcal{A} = 18$	—	—

Input: $(B, 4, 84, 84)$ float32 in $[0, 1]$. After the two conv layers the spatial size is 9×9 , yielding $32 \times 81 = 2592$ flattened features.

4 Final Hyperparameters

The final configuration used for the three reported runs is:

Hyperparameter	Value
Optimizer	RMSprop
Learning rate	0.00025
RMSprop α (decay)	0.95
RMSprop ε	0.01
Discount factor γ	0.97
Batch size	64
Replay buffer capacity	1,000,000
Min replay size before training	50,000
Update frequency	every 4 env steps
Loss function	MSE
Gradient clipping	none
ε -greedy start	1.0
ε -greedy end	0.001
ε anneal steps	250,000
ε at evaluation	0.001
Frame skip	4
Frame stack	4
Life-loss terminal	true
No-op max	0 (disabled)
Total training steps	5,000,000
LR decay start step	3,000,000
LR decay end (lr_min)	0.0
Best-checkpoint final eval (post-training)	true
Mid-training eval frequency	every 10,000 steps
Mid-training eval episodes	1
Final evaluation episodes	100

Key deviations from paper defaults. Three hyperparameters are deliberately tuned away from the 2013 paper defaults:

- $\gamma = 0.97$ instead of 0.99 — shorter bootstrap horizon accelerates peak convergence.
- LR warmdown: linear decay of the learning rate from 0.00025 to 0 starting at step 3M.
- Best-checkpoint evaluation: the final 100-episode result is evaluated at the peak-performance checkpoint rather than the final (post-collapse) checkpoint.

5 Training Protocol

- **Training steps:** 5,000,000 environment steps per run. A single call to `env.step()` counts as one step; with frame-skip=4, this corresponds to approximately 1.25M agent actions.

- **Mid-training evaluation:** every 10,000 steps, one full episode is played with $\varepsilon = 0.001$ (near-greedy). The episode reward is logged to `training_stats.jsonl`. These single-episode evaluations are used only for tracking and checkpoint selection — they are *not* the reported result.
- **Seeds:** three independent runs using seeds 1, 2, and 3.

Final Evaluation Protocol (Post-Training)

The final reported score is produced *after* training ends, as a separate evaluation step:

1. During training, whenever a new single-episode evaluation high is reached, the current network weights are saved as a **best checkpoint** (`ckpt_best.pt`). This is a bookkeeping operation only — it does not modify the training procedure in any way.
2. After training completes, the best checkpoint is loaded and **100 full episodes** are played with $\varepsilon = 0.001$. The mean reward over these 100 episodes is the final reported score.

This decouples the *evaluation point* from the end of training, which matters because the single-network DQN (without a target network) suffers from post-peak collapse: the weights at the end of 5M steps are degraded relative to the peak. The best-checkpoint approach captures the peak-quality policy while still training for the full budget.

6 Training Curves

Figure 1 shows the mid-training evaluation reward (1-episode, noisy) over the full 5M training steps for all three seeds, overlaid on the same axes. Figure 2 shows each seed individually with the best-checkpoint step marked.




Figure 1: Evaluation reward (1 episode, every 10k steps) for seeds 1, 2, 3. Thin lines are raw; thick lines are a 20-point rolling mean. Dashed vertical lines mark the best-checkpoint step selected for the final 100-episode evaluation.




Figure 2: Per-seed training curves. Red dashed vertical line: step of best checkpoint used for final eval. Red dotted horizontal line: final 100-episode score.

7 Final Results

Run	Best ckpt step	Final reward (100 ep)	Notes
Seed 1	3,630,000	3840.4	Highest single-seed score
Seed 2	4,150,000	2113.1	Solid, near-greedy exploitation
Seed 3	1,920,000	3558.3	Fast convergence, early peak
Average	—	3170.6	Bonus threshold: 3000 ✓

8 What Worked

We ran 58 experiments across roughly 800 GPU-hours, systematically searching for the configuration that maximises the 100-episode final evaluation reward. The following changes each contributed positively:

8.1 Life-loss terminal signal

Treating each in-game life loss as a terminal state for the Q-value bootstrap improved performance consistently across all seeds and learning rates. It gives the agent a direct incentive to survive, shaping rewards more densely.

8.2 Fast epsilon annealing (250k steps)

Annealing ϵ from 1.0 to 0.001 over only 250k steps (rather than 1M as in the paper) was the key to achieving “sweet-zone” peak timing: the agent commits to near-greedy exploitation early, and Q-values peak in the 2–4M step range. Longer annealing windows (500k, 1M) led to bimodal and worse results.

8.3 Lower discount factor ($\gamma = 0.97$)

Reducing γ from 0.99 to 0.97 shortens the bootstrap horizon. The single-network DQN without a target network amplifies Q-values over time. With $\gamma = 0.97$, the policy converges faster and peaks higher (single-episode peaks of 4000–4260 vs. 2000 with $\gamma = 0.99$).

8.4 LR warmdown (lr_decay_start_step = 3M)

After step 3M the learning rate is linearly decayed from 0.00025 to 0 over the remaining 2M steps. This “fine-tuning” phase slows the Q-value divergence that causes post-peak collapse and stabilises the policy near its peak. Without this, all seeds collapsed sharply after their peak.

8.5 Best-checkpoint evaluation (use_best_checkpoint_eval)

This was the single most impactful change in the entire search, improving average reward from 427 to 1627 in one step. During training, `ckpt_best.pt` is overwritten whenever a new 1-episode evaluation high is reached. At the end of training the 100-episode final evaluation is run on this best-checkpoint rather than the degraded final weights.

Motivation: the root cause of poor performance was not slow learning — the agent consistently reached excellent policies mid-training (peaks of 1500–4260 in single-episode evals) — but the single-network bootstrap loop caused catastrophic policy collapse in late training. Peak reward and final reward were nearly uncorrelated without this change.

8.6 Stable learning-rate regime

We found that `lr = 0.00025` with `batch_size=64` is the only stable, high-performing combination. Smaller batches (32) caused explosive Q-value divergence; larger learning rates destabilised training from step 0.

9 What Did Not Work

The following configurations were tested and conclusively ruled out:

Change	Avg	Why it failed
$\gamma = 0.97$ <i>without</i> best-ckpt	338.9	Faster collapse erases high peaks before eval
Huber loss	182.3	Clips early gradients; agent barely learns, peaks at step 460k
lr = 0.0001	276.0	Peaks too late (step 4.8M+); collapse hits exactly at deadline
lr = 0.0001 + noop + life	400.0	Same late-peak problem
update_freq=2	268.7	2× more updates/step → Q-values diverge faster
update_freq=8	1233.7	Fewer updates → agent learns too slowly, lower peaks
eps_end=0.005	367.7	Extra noise kills lucky seeds (s2: 825→346)
eps_anneal=500k	417.6	Bimodal peak timing; many seeds peak too early or too late
rmsprop_eps=0.1	226.7	Uniformly smaller updates → faster collapse (677/Mstep)
noop_max=30	consistent hurt	Adds noise to start states; consistently worse in 4 exps
min_replay_size=100k	1730.0	Delays first update, reducing time at full lr; hurts peaks
grad_clip + lr = 0.00025	diverge	Clipping disrupts RMSprop’s variance accumulator
frame_differencing	disallowed	Not in 2013 paper

9.1 The post-peak collapse problem

The central challenge throughout was that the single-network DQN (without a target network) suffers from bootstrap instability: the same network is used both to generate targets and to be updated, creating a self-referential loop that amplifies Q-values until the policy becomes incoherent. Every “within-training” attempt to slow this collapse failed:

- Lower γ *accelerated* collapse (shorter horizon means faster self-amplification per step).
- Huber loss slowed early learning too much.
- Lower lr delayed the peak past the training deadline.
- rmsprop_eps= 0.1 made updates uniformly smaller, not selectively smaller in late training.

The solution was orthogonal: instead of preventing collapse, we *decouple the evaluation point from the end of training* via best-checkpoint evaluation, capturing the policy at peak quality regardless of subsequent degradation.

10 Implementation Notes

- **No RL library used.** All components (replay buffer, agent, preprocessing, evaluation loop) are implemented from scratch in PyTorch.
- **AMP (mixed precision):** training uses `torch.cuda.amp` for speed.
- **Checkpointing:** model weights and optimizer state are saved to `runs/<run_id>/ckpt_best.pt` and `ckpt_NNNNNNNNNN.pt` every 500k steps.
- **Logging:** all training metrics (loss, Q-values, epsilon, lr, eval reward) are logged to Weights & Biases and to `training_stats.jsonl` for offline analysis.
- **Compute:** training was run on NVIDIA GPUs via SLURM.

11 Conclusion

Starting from a baseline DQN on Alien scoring roughly 427, we systematically identified and resolved the root cause of poor performance: post-peak policy collapse caused by single-network bootstrap instability. The key insight was that the agent consistently learned excellent policies mid-training — the obstacle was entirely in the *evaluation protocol*.

By combining:

1. $\gamma = 0.97$ (higher, faster-converging peaks),
2. linear LR warmdown from step 3M (slower collapse in the fine-tuning phase), and
3. best-checkpoint evaluation (100-episode final eval at the peak, not the collapsed end),

we achieved a 3-run average of **3170.6**, exceeding the DQN paper bonus target of 3000.